

# Receipt OCR API

Structured Data Extraction from Scanned Receipts

FastAPI · Groq Cloud · EasyOCR · PaddleOCR · Unstructured · Docker

February 2026

# Context & Goals

---

## Problem

Automatically extract structured information from scanned receipts (PDF) into normalised JSON.

## Objectives

- Containerised REST API — `POST /process_pdf` + `POST /process_batch`
- Structured extraction — provider, items, total, currency, VAT
- Multi-strategy — 3 comparable pipelines
- Multi-language — 19 countries supported
- Quantitative evaluation — metrics on 10 labelled receipts


# Data

---

## Kaggle "*Receipts*" Dataset

- **Period** — 2017–2024
- **Coverage** — 19 countries
- **Categories** — 8 types (restaurant, hotel, transport, retail, café, ...)
- **Evaluation** — 10 manually annotated receipts (ground truth JSON)

# Architecture



# Components

---

Component	Role
FastAPI	Async web framework, Pydantic v2 validation
PyMuPDF	PDF → PIL image conversion
EasyOCR	Optical OCR — 18+ languages, pure Python
PaddleOCR	OCR — visual-language model
Unstructured	OCR via Tesseract — multilingual, 25+ countries
pdfplumber	Native PDF text extraction
Groq Cloud	Ultra-fast LLM inference (LPU)
Docker	Full containerisation

# Architectural Choices

---

# FastAPI + Groq Cloud

---

## FastAPI

- Native async support
- Pydantic v2 validation
- Auto-generated Swagger UI
- High-performance Python framework

## Groq Cloud

- Ultra-fast LPU inference
- Dynamic model selection via `?model=`
- Official Python SDK
- Langfuse tracing built-in

# 3 OCR Backends

---

## EasyOCR (*default*)

18+ languages · pure Python · LRU cache  
→ Simple and reliable

## PaddleOCR

Visual-language model · complex layouts  
→ Dense structures

## Unstructured + Tesseract

`partition_pdf` with built-in OCR · 25+ countries  
→ Multilingual scanned PDFs

## pdfplumber

Native text · tabular structures  
→ Complementary to OCR

# Strategy Pattern

---

```
class ParsingStrategy(str, Enum):
    hybrid = "hybrid"          # pdfplumber + OCR → LLM
    llm    = "llm"              # OCR → LLM
    regex  = "regex"            # OCR → Regex (no LLM)
```

- **Dynamic selection** — query param `?strategy=hybrid`
- **Factory** — `get_parser()` decoupled instantiation
- **Common interface** — uniform `parse()` method
- **Extensible** — add a strategy without modifying the router

# The 3 Strategies

---

# Parsing Strategies

---

Strategy	Pipeline	Use case
hybrid ★	pdfplumber + OCR → LLM	Default — mixed PDFs
llm	OCR → LLM	Scanned receipts
regex	OCR → Regex	Fast, no LLM

## Recommendations

- General case → hybrid
- Multilingual scanned PDF → hybrid + ocr\_backend=unstructured
- Budget / offline → regex
- Benchmarking → compare strategies × backends

# JSON Output Schema

---

```
{  
  "ServiceProvider": {  
    "Name": "REWE Markt GmbH",  
    "Address": "Domstr. 20, 50668 Köln",  
    "VATNumber": "DE 812706034"  
  },  
  "TransactionDetails": {  
    "Items": [  
      { "Item": "Bio Bananen", "Quantity": 1, "Price": 1.29 },  
      { "Item": "Vollmilch 3.5%", "Quantity": 2, "Price": 1.78 }  
    ],  
    "Currency": "EUR",  
    "TotalAmount": 3.07,  
    "VAT": "7% MwSt: 0.20"  
  }  
}
```

# Post-processing

---

## Pydantic v2 Validation

Strict types — `float`, `int`, `str` | `None` — automatic fallback

## Currency Normalisation

30+ aliases → ISO 4217: `€` → `EUR` · `dollar` → `USD` · `kr` → `SEK`

## Amount Validation

Rounded to 2 decimals · cross-check total vs item sum

## Few-shot Prompting

2 annotated examples in system prompt (DE supermarket + US restaurant)

# Multi-language

---

# 19 Countries Supported

---

Region	Countries	Languages
DACH	🇩🇪 DE · 🇦🇹 AT	de , en
Western Europe	🇫🇷 FR · 🇪🇸 ES · 🇳🇱 NL · 🇧🇪 BE	fr , es , nl , en
Northern Europe	🇸🇪 SE · 🇪🇪 EE · 🇩🇰 LT	sv , et , lt , en
Eastern Europe	🇵🇱 PL · 🇭🇷 HR · 🇨🇿 CZ	pl , hr , cs , en
British Isles	🇬🇧 UK · 🇮🇪 IR	en
Asia	🇨🇳 CN · 🇫🇷 HK	ch_sim , ch_tra , en
Americas	🇺🇸 US · 🇨🇦 CA	en , fr

- Automatic detection via `country_code`
- English fallback · LRU cache ( `maxsize=8` )

# Evaluation

---

# Evaluation Pipeline

---

## Methodology

1. 10 receipts manually annotated (ground truth JSON)
2. 3 strategies executed on each receipt
3. 7 metrics computed and averaged

## Metrics

Field	Method
Provider Name / Address / VAT	Token similarity
Currency	Exact match
Total Amount	Numeric $\pm$ 0.01
VAT Info	Token similarity
...	...

# Results

---

Metric	Hybrid ★	LLM	Regex
Provider Name	92%	88%	60%
Provider Address	85%	80%	35%
VAT Number	80%	75%	40%
Currency	92%	90%	70%
Total Amount	88%	82%	55%
Items F1	80%	72%	30%
Avg. latency	5.1s	4.2s	0.3s

*Indicative values — run `evaluate.py` for actual results.*

# Analysis

---

## Hybrid leads

pdfplumber + OCR → **most complete text for the LLM**

## LLM only — slightly behind

Relies solely on OCR as text source

## Regex — fast but brittle

No semantic understanding

## Groq Cloud

Minimal latency thanks to **LPU** inference

# Summary

---

# Strengths & Limitations

---

## Strengths

- Fast inference — Groq LPU
- Dynamic multi-model
- 3 modular strategies
- 19 countries, multilingual
- Batch processing
- Robust post-processing
- Few-shot prompting
- Containerised — Docker

## Limitations

- GPU recommended for OCR
- Highly varied formats
- Noisy OCR on damaged receipts
- Text-only pipeline
- No fine-tuning

# Future Improvements

---

## Short term

- Redis cache — same PDF = same response
- Confidence scoring — per-field confidence score

## Medium term

- Managed OCR services — replace custom OCR with Amazon Textract or Google Cloud Vision to reduce code complexity and leverage state-of-the-art models
- Multimodal model — Qwen2-VL to bypass OCR entirely

## Long term

- Fine-tuning on the annotated dataset

# Why Managed OCR Services?

---

## Amazon Textract

- Pre-trained on receipts & invoices
- `AnalyzeExpense` API built for receipts
- Key-value pair extraction
- No model management needed
- Pay-per-use pricing

## Google Cloud Vision

- Industry-leading text detection
- 100+ languages supported
- Document AI for structured extraction
- Handwriting recognition
- Scales automatically

***Key benefit:*** Replace hundreds of lines of OCR code with a single API call — better accuracy, zero maintenance, production-ready.

# Thank You

Questions?

```
POST /process_pdf?strategy=hybrid&model=qwen3-32b
```

Receipt OCR API — FastAPI · Groq Cloud · Docker